

```

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from keras.models import Model
from keras.layers import LSTM, Activation, Dense, Dropout,
Input, Embedding
from keras.optimizers import RMSprop
from keras.preprocessing.text import Tokenizer
from keras_preprocessing import sequence
from keras.utils import to_categorical
from keras.models import load_model
import csv
import tensorflow as tf
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
import nltk
nltk.download('stopwords')
from nltk.corpus import stopwords
STOPWORDS = set(stopwords.words('english'))

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Unzipping corpora/stopwords.zip.

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
df = pd.read_csv('/content/spam.csv', delimiter=',', encoding='latin-1')
df.head()

```

	v1	v2	Unnamed: 2
0	ham	Go until jurong point, crazy.. Available only ...	NaN
1	ham	Ok lar... Joking wif u oni...	NaN
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	NaN
3	ham	U dun say so early hor... U c already then say...	NaN
4	ham	Nah I don't think he goes to usf, he lives aro...	NaN

```

      Unnamed: 3 Unnamed: 4
0      NaN      NaN

```

1	NaN	NaN
2	NaN	NaN
3	NaN	NaN
4	NaN	NaN

```
df.drop(['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'], axis=1, inplace=True)
#dropping unwanted columns
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5572 entries, 0 to 5571
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype
---  ---
 0    v1     5572 non-null    object
 1    v2     5572 non-null    object
dtypes: object(2)
memory usage: 87.2+ KB
```

```
#Count of Spam and Ham values
df.groupby(['v1']).size()
```

```
v1
ham      4825
spam      747
dtype: int64
```

```
#Label Encoding target column
```

```
X = df.v2
Y = df.v1
le = LabelEncoder()
Y = le.fit_transform(Y)
Y = Y.reshape(-1,1)
```

```
# Test and train split
```

```
X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.15)
```

```
# Tokenisation function
```

```
max_words = 1000
max_len = 150
tok = Tokenizer(num_words=max_words)
tok.fit_on_texts(X_train)
sequences = tok.texts_to_sequences(X_train)
sequences_matrix = sequence.pad_sequences(sequences,maxlen=max_len)
```

```
#creating LSTM model
```

```
inputs = Input(name='InputLayer',shape=[max_len])
layer = Embedding(max_words,50,input_length=max_len)(inputs)
layer = LSTM(64)(layer)
layer = Dense(256,name='FullyConnectedLayer1')(layer)
layer = Activation('relu')(layer)
layer = Dropout(0.5)(layer)
```

```

layer = Dense(1,name='OutputLayer')(layer)
layer = Activation('sigmoid')(layer)

model = Model(inputs=inputs,outputs=layer)
model.summary()
model.compile(loss='binary_crossentropy',optimizer=RMSprop(),metrics=[
'accuracy'])

```

Model: "model"

Layer (type)	Output Shape	Param #
InputLayer (InputLayer)	[(None, 150)]	0
embedding (Embedding)	(None, 150, 50)	50000
lstm (LSTM)	(None, 64)	29440
FullyConnectedLayer1 (Dense)	(None, 256)	16640
activation (Activation)	(None, 256)	0
dropout (Dropout)	(None, 256)	0
OutputLayer (Dense)	(None, 1)	257
activation_1 (Activation)	(None, 1)	0

```

Total params: 96,337
Trainable params: 96,337
Non-trainable params: 0

```

```

model.fit(sequences_matrix,Y_train,batch_size=128,epochs=10,
validation_split=0.2)

```

```

Epoch 1/10
30/30 [=====] - 11s 278ms/step - loss: 0.3272
- accuracy: 0.8712 - val_loss: 0.1428 - val_accuracy: 0.9589
Epoch 2/10
30/30 [=====] - 8s 259ms/step - loss: 0.0858
- accuracy: 0.9789 - val_loss: 0.0664 - val_accuracy: 0.9800
Epoch 3/10
30/30 [=====] - 8s 259ms/step - loss: 0.0429
- accuracy: 0.9892 - val_loss: 0.0830 - val_accuracy: 0.9789
Epoch 4/10
30/30 [=====] - 8s 259ms/step - loss: 0.0335
- accuracy: 0.9910 - val_loss: 0.0504 - val_accuracy: 0.9842
Epoch 5/10

```

```

30/30 [=====] - 12s 392ms/step - loss: 0.0227
- accuracy: 0.9934 - val_loss: 0.0547 - val_accuracy: 0.9852
Epoch 6/10
30/30 [=====] - 8s 259ms/step - loss: 0.0192
- accuracy: 0.9939 - val_loss: 0.0571 - val_accuracy: 0.9831
Epoch 7/10
30/30 [=====] - 10s 332ms/step - loss: 0.0112
- accuracy: 0.9971 - val_loss: 0.1148 - val_accuracy: 0.9778
Epoch 8/10
30/30 [=====] - 8s 258ms/step - loss: 0.0108
- accuracy: 0.9974 - val_loss: 0.0783 - val_accuracy: 0.9831
Epoch 9/10
30/30 [=====] - 8s 259ms/step - loss: 0.0081
- accuracy: 0.9976 - val_loss: 0.0760 - val_accuracy: 0.9842
Epoch 10/10
30/30 [=====] - 8s 260ms/step - loss: 0.0052
- accuracy: 0.9989 - val_loss: 0.0963 - val_accuracy: 0.9821

```

<keras.callbacks.History at 0x7f8b53cd1d10>

```
model.save("model_1")
```

```

WARNING:absl:Function `_wrapped_model` contains input name(s)
InputLayer with unsupported characters which will be renamed to
inputlayer in the SavedModel.
WARNING:absl:Found untraced functions such as lstm_cell_layer_call_fn,
lstm_cell_layer_call_and_return_conditional_losses while saving
(showing 2 of 2). These functions will not be directly callable after
loading.

```

```

test_sequences = tok.texts_to_sequences(X_test)
test_sequences_matrix
=sequence.pad_sequences(test_sequences,maxlen=max_len)
accuracy = model.evaluate(test_sequences_matrix,Y_test)
print('Accuracy: {:.3f}'.format(accuracy[1]))

```

```

27/27 [=====] - 1s 22ms/step - loss: 0.1190 -
accuracy: 0.9773
Accuracy: 0.977

```

```

y_pred = model.predict(test_sequences_matrix)
print(y_pred[25:40].round(3))

```

```

27/27 [=====] - 2s 41ms/step
[[0.   ]
 [0.   ]
 [0.   ]
 [0.003]
 [0.   ]
 [0.   ]
 [0.   ]
 [0.   ]

```

```
[0.  ]
[0.  ]
[0.  ]
[0.002]
[0.  ]
[0.  ]
[1.  ]]
```

```
print(Y_test[25:40])
```

$$\begin{bmatrix} [0] \\ [0] \\ [0] \\ [0] \\ [0] \\ [0] \\ [0] \\ [0] \\ [0] \\ [1] \\ [0] \\ [0] \\ [0] \\ [1] \end{bmatrix}$$